


Lecture 10-1

ReLU: Better non-linearity


Sung Kim <hunkim+mr@gmail.com>
<http://hunkim.github.io/ml/>

NN for XOR



NN for XOR

```
[ ] W1 = tf.Variable(tf.random_uniform([2, 2], -1.0, 1.0))  
[ ] W2 = tf.Variable(tf.random_uniform([2, 1], -1.0, 1.0))  
  
[ ] b1 = tf.Variable(tf.zeros([2]), name="Bias1")  
[ ] b2 = tf.Variable(tf.zeros([1]), name="Bias2")  
  
# Our hypothesis  
L2 = tf.sigmoid(tf.matmul(X, W1) + b1)  
hypothesis = tf.sigmoid(tf.matmul(L2, W2) + b2)
```



Let's go deep & wide!

```

W1 = tf.Variable(tf.random_uniform([2, 5], -1.0, 1.0))
W2 = tf.Variable(tf.random_uniform([5, 4], -1.0, 1.0))
W3 = tf.Variable(tf.random_uniform([4, 1], -1.0, 1.0))

```

```

b1 = tf.Variable(tf.zeros([5]), name="Bias1")
b2 = tf.Variable(tf.zeros([4]), name="Bias2")
b3 = tf.Variable(tf.zeros([1]), name="Bias2")

```

Our hypothesis

```

L2 = tf.sigmoid(tf.matmul(X, W1) + b1)
L3 = tf.sigmoid(tf.matmul(L2, W2) + b2)
hypothesis = tf.sigmoid(tf.matmul(L3, W3) + b3)


```

$X \{$

x_1, x_2

\downarrow

$bias$



9 hidden layers!

```
W1 = tf.Variable(tf.random_uniform([2, 5], -1.0, 1.0), name = "Weight1")
```

```
W2 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight2")
```

```
W3 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight3")
```

```
W4 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight4")
```

```
W5 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight5")
```

```
W6 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight6")
```

```
W7 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight7")
```

```
W8 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight8")
```

```
W9 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight9")
```

```
W10 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight10")
```

```
W11 = tf.Variable(tf.random_uniform([5, 1], -1.0, 1.0), name = "Weight11")
```

```
b1 = tf.Variable(tf.zeros([5]), name="Bias1")
```

```
b2 = tf.Variable(tf.zeros([5]), name="Bias2")
```

```
b3 = tf.Variable(tf.zeros([5]), name="Bias3")
```

```
b4 = tf.Variable(tf.zeros([5]), name="Bias4")
```

```
b5 = tf.Variable(tf.zeros([5]), name="Bias5")
```

```
b6 = tf.Variable(tf.zeros([5]), name="Bias6")
```

```
b7 = tf.Variable(tf.zeros([5]), name="Bias7")
```

```
b8 = tf.Variable(tf.zeros([5]), name="Bias8")
```

```
b9 = tf.Variable(tf.zeros([5]), name="Bias9")
```

```
b10 = tf.Variable(tf.zeros([5]), name="Bias10")
```

```
b11 = tf.Variable(tf.zeros([1]), name="Bias11")
```

9 hidden layers!

```
W1 = tf.Variable(tf.random_uniform([2, 5], -1.0, 1.0), name = "Weight1")
W2 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight2")
W3 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight3")
W4 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight4")
W5 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight5")
W6 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight6")
W7 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight7")
W8 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight8")
W9 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight9")
W10 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight10")

W11 = tf.Variable(tf.random_uniform([5, 1], -1.0, 1.0), name = "Weight11")


b1 = tf.Variable(tf.zeros([5]), name="Bias1")
b2 = tf.Variable(tf.zeros([5]), name="Bias2")
b3 = tf.Variable(tf.zeros([5]), name="Bias3")
b4 = tf.Variable(tf.zeros([5]), name="Bias4")
b5 = tf.Variable(tf.zeros([5]), name="Bias5")
b6 = tf.Variable(tf.zeros([5]), name="Bias6")
b7 = tf.Variable(tf.zeros([5]), name="Bias7")
b8 = tf.Variable(tf.zeros([5]), name="Bias8")
b9 = tf.Variable(tf.zeros([5]), name="Bias9")
b10 = tf.Variable(tf.zeros([5]), name="Bias10")

b11 = tf.Variable(tf.zeros([1]), name="Bias11")
```

Our hypothesis

```
L1 = tf.sigmoid(tf.matmul(X, W1) + b1)
L2 = tf.sigmoid(tf.matmul(L1, W2) + b2)
L3 = tf.sigmoid(tf.matmul(L2, W3) + b3)
L4 = tf.sigmoid(tf.matmul(L3, W4) + b4)
L5 = tf.sigmoid(tf.matmul(L4, W5) + b5)
L6 = tf.sigmoid(tf.matmul(L5, W6) + b6)
L7 = tf.sigmoid(tf.matmul(L6, W7) + b7)
L8 = tf.sigmoid(tf.matmul(L7, W8) + b8)
L9 = tf.sigmoid(tf.matmul(L8, W9) + b9)
L10 = tf.sigmoid(tf.matmul(L9, W10) + b10)
```

hypothesis = tf.sigmoid(tf.matmul(L10, W11) + b11)



9 hidden layers!

```
W1 = tf.Variable(tf.random_uniform([2, 5], -1.0, 1.0), name = "Weight1")
W2 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight2")
W3 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight3")
W4 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight4")
W5 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight5")
W6 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight6")
W7 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight7")
W8 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight8")
W9 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight9")
W10 = tf.Variable(tf.random_uniform([5, 5], -1.0, 1.0), name = "Weight10")

W11 = tf.Variable(tf.random_uniform([5, 1], -1.0, 1.0), name = "Weight11")


b1 = tf.Variable(tf.zeros([5]), name="Bias1")
b2 = tf.Variable(tf.zeros([5]), name="Bias2")
b3 = tf.Variable(tf.zeros([5]), name="Bias3")
b4 = tf.Variable(tf.zeros([5]), name="Bias4")
b5 = tf.Variable(tf.zeros([5]), name="Bias5")
b6 = tf.Variable(tf.zeros([5]), name="Bias6")
b7 = tf.Variable(tf.zeros([5]), name="Bias7")
b8 = tf.Variable(tf.zeros([5]), name="Bias8")
b9 = tf.Variable(tf.zeros([5]), name="Bias9")
b10 = tf.Variable(tf.zeros([5]), name="Bias10")

b11 = tf.Variable(tf.zeros([1]), name="Bias11")
```

```
# Our hypothesis
with tf.name_scope("layer1") as scope:
    L1 = tf.sigmoid(tf.matmul(X, W1) + b1)
with tf.name_scope("layer2") as scope:
    L2 = tf.sigmoid(tf.matmul(L1, W2) + b2)
with tf.name_scope("layer3") as scope:
    L3 = tf.sigmoid(tf.matmul(L2, W3) + b3)
with tf.name_scope("layer4") as scope:
    L4 = tf.sigmoid(tf.matmul(L3, W4) + b4)
with tf.name_scope("layer5") as scope:
    L5 = tf.sigmoid(tf.matmul(L4, W5) + b5)
with tf.name_scope("layer6") as scope:
    L6 = tf.sigmoid(tf.matmul(L5, W6) + b6)
with tf.name_scope("layer7") as scope:
    L7 = tf.sigmoid(tf.matmul(L6, W7) + b7)
with tf.name_scope("layer8") as scope:
    L8 = tf.sigmoid(tf.matmul(L7, W8) + b8)
with tf.name_scope("layer9") as scope:
    L9 = tf.sigmoid(tf.matmul(L8, W9) + b9)
with tf.name_scope("layer10") as scope:
    L10 = tf.sigmoid(tf.matmul(L9, W10) + b10)


with tf.name_scope("last") as scope:
    hypothesis = tf.sigmoid(tf.matmul(L10, W11) + b11)
```

Tensorboard visualization




Poor results?


```
196000 [0.69314718, array([[ 0.49999988],  
[ 0.50000006],  
[ 0.49999982],  
[ 0.5        ]], dtype=float32)]  
198000 [0.69314718, array([[ 0.49999988],  
[ 0.50000006],  
[ 0.49999982],  
[ 0.5        ]], dtype=float32)]  
[array([[ 0.49999988],  
[ 0.50000006],  
[ 0.49999982],  
[ 0.5        ]], dtype=float32), array([[ 0.],  
[ 1.],  
[ 0.],  
[ 1.]], dtype=float32)]  
Accuracy: 0.5
```



cost




accuracy




Tensorboard
Cost &
Accuracy

Backpropagation


Q6



lec 9-2: Backpropagation (chain rule)



Vanishing gradient (NN winter2: 1986-2006)




Geoffrey Hinton's summary of findings up to today


- Our labeled datasets were thousands of times too small.
- Our computers were millions of times too slow.
- We initialized the weights in a stupid way.
- We used the wrong type of non-linearity.




Sigmoid!



Sigmoid!




ReLU: Rectified Linear Unit



ReLU: Rectified Linear Unit

$$\max(0, x)$$



$L1 = \underline{\text{tf.sigmoid}}(\text{tf.matmul}(X, W1) + b1)$

$L1 = \boxed{\text{tf.nn.relu}}(\text{tf.matmul}(X, W1) + b1)$

ReLU




```
# Our hypothesis
with tf.name_scope("layer1") as scope:
    L1 = tf.nn.relu(tf.matmul(X, W1) + b1)
with tf.name_scope("layer2") as scope:
    L2 = tf.nn.relu(tf.matmul(L1, W2) + b2)
with tf.name_scope("layer3") as scope:
    L3 = tf.nn.relu(tf.matmul(L2, W3) + b3)
with tf.name_scope("layer4") as scope:
    L4 = tf.nn.relu(tf.matmul(L3, W4) + b4)
with tf.name_scope("layer5") as scope:
    L5 = tf.nn.relu(tf.matmul(L4, W5) + b5)
with tf.name_scope("layer6") as scope:
    L6 = tf.nn.relu(tf.matmul(L5, W6) + b6)
with tf.name_scope("layer7") as scope:
    L7 = tf.nn.relu(tf.matmul(L6, W7) + b7)
with tf.name_scope("layer8") as scope:
    L8 = tf.nn.relu(tf.matmul(L7, W8) + b8)
with tf.name_scope("layer9") as scope:
    L9 = tf.nn.relu(tf.matmul(L8, W9) + b9)
with tf.name_scope("layer10") as scope:
    L10 = tf.nn.relu(tf.matmul(L9, W10) + b10)

with tf.name_scope("last") as scope:
    hypothesis = tf.sigmoid(tf.matmul(L10, W11) + b11)
```

Works very well


```
196000 [2.6226094e-06, array([[ 2.59195826e-06],  
[ 9.99999642e-01],  
[ 9.99994874e-01],  
[ 2.43454133e-06]], dtype=float32)]  
198000 [2.607708e-06, array([[ 2.55822852e-06],  
[ 9.99999642e-01],  
[ 9.99994874e-01],  
[ 2.40260101e-06]], dtype=float32)]  
[array([[ 2.52509381e-06],  
[ 9.99999642e-01],  
[ 9.99994874e-01],  
[ 2.37124709e-06]], dtype=float32), array([[ 0.],  
[ 1.],  
[ 1.],  
[ 0.]], dtype=float32)]  
Accuracy: 1.0
```

Works very well



Cost function


cost



Activation Functions


Sigmoid

$$\sigma(x) = 1/(1 + e^{-x})$$




tanh

tanh(x)




ReLU

$\max(0, x)$



~~Leaky ReLU~~
 $\max(0.1x, x)$




Maxout

ELU

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha (\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$



Activation functions on CIFAR-10

	maxout	ReLU	VLReLU	tanh	Sigmoid
	93.94	92.11	92.97	89.28	n/c
	93.78	91.74	92.40	89.48	n/c
-	—	91.93	93.09	—	n/c
	91.75	90.63	92.27	89.82	n/c
	n/c†	90.91	92.43	89.54	n/c

[Mishkin et al. 2015]

Next
Weight initialization




Lecture 10-2

Initialize weights in a smart way

Sung Kim <hunkim+mr@gmail.com>
<http://hunkim.github.io/ml/>



Vanishing gradient



Geoffrey Hinton's summary of findings up to today


- Our labeled datasets were thousands of times too small.
- Our computers were millions of times too slow.
- We initialized the weights in a stupid way.
- We used the wrong type of non-linearity.

Set all initial weights to 0



Cost function


cost



Need to set the initial weight values wisely

- Not all 0's
- Challenging issue
- Hinton et al. (2006) "A Fast Learning Algorithm for Deep Belief Nets"
 - Restricted Boatman Machine (RBM)


RBM STRUCTURE



RESTRICTION = NO CONNECTIONS
WITHIN A LAYER

RECREATE INPUT


Forward



Backward




RECREATE INPUT



RECREATE INPUT


Forward




Backward

RECREATE INPUT

Forward




Backward




RECREATE INPUT

Forward




Backward




RECREATE INPUT

Forward




Backward




RECREATE INPUT


Forward



Backward




RECREATE INPUT




KL DIVERGENCE = COMPARE
ACTUAL TO RECREATION


How can we use RBM to initialize weights?




- Apply the RBM idea on adjacent two layers as a pre-training step
- Continue the first process to all layers
- This will set weights
- Example: Deep Belief Network
 - Weight initialized by RBM




PRE-TRAINING




PRE-TRAINING







PRE-TRAINING



PRE-TRAINING



FINE TUNING



Good news



- No need to use complicated RBM for weight initializations
- Simple methods are OK
 - Xavier initialization: X. Glorot and Y. Bengio, “Understanding the difficulty of training deep feedforward neural networks,” in International conference on artificial intelligence and statistics, 2010
 - He’s initialization: K. He, X. Zhang, S. Ren, and J. Sun, “Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification,” 2015

Xavier/He initialization

- Makes sure the weights are ‘just right’, not too small, not too big
- Using number of input (fan_in) and output (fan_out)

```
# Xavier initialization
# Glorot et al. 2010
W = np.random.randn(fan_in, fan_out)/np.sqrt(fan_in)
```

He et al. 2015

```
W = np.random.randn(fan_in, fan_out)/np.sqrt(fan_in/2)
```

prettytensor implementation


```
def xavier_init(n_inputs, n_outputs, uniform=True):
    """Set the parameter initialization using the method described.
    This method is designed to keep the scale of the gradients roughly the same
    in all layers.
    Xavier Glorot and Yoshua Bengio (2010):
        Understanding the difficulty of training deep feedforward neural
        networks. International conference on artificial intelligence and
        statistics.

    Args:
        n_inputs: The number of input nodes into each output.
        n_outputs: The number of output nodes for each input.
        uniform: If true use a uniform distribution, otherwise use a normal.

    Returns:
        An initializer.
    """
    if uniform:
        # 6 was used in the paper.
        init_range = math.sqrt(6.0 / (n_inputs + n_outputs))
        return tf.random_uniform_initializer(-init_range, init_range)
    else:
        # 3 gives us approximately the same limits as above since this repicks
        # values greater than 2 standard deviations from the mean.
        stddev = math.sqrt(3.0 / (n_inputs + n_outputs))
        return tf.truncated_normal_initializer(stddev=stddev)
```

<http://stackoverflow.com/questions/33640581/how-to-do-xavier-initialization-on-tensorflow>

Activation functions and initialization on CIFAR-10



Init method	maxout	ReLU	VLReLU	tanh	Sigmoid
LSUV	93.94	92.11	92.97	89.28	n/c
OrthoNorm	93.78	91.74	92.40	89.48	n/c
OrthoNorm-MSRA scaled	–	91.93	93.09	–	n/c
Xavier	91.75	90.63	92.27	89.82	n/c
MSRA	n/c†	90.91	92.43	89.54	n/c

[Mishkin et al. 2015]

Still an active area of research


- We don't know how to initialize perfect weight values, yet
- Many new algorithms
 - Batch normalization
 - Layer sequential uniform variance
 - ...



Geoffrey Hinton's summary of findings up to today

- Our labeled datasets were thousands of times too small.
- Our computers were millions of times too slow.
- We initialized the weights in a stupid way. ↴
- We used the wrong type of non-linearity. ↴

Next
dropout and
model ensemble





Lecture 10-3


NN dropout and model ensemble

Sung Kim <hunkim+mr@gmail.com>

Overfitting




Am I overfitting?




- Very high accuracy on the training dataset (eg: 0.99)
- Poor accuracy on the test data set (0.85)

Solutions for overfitting

- More training data!
 - Reduce the number of features
 - Regularization
- 


Regularization

- Let's not have too big numbers in the weight




Regularization

- Let's not have too big numbers in the weight




```
l2reg = 0.001 * tf.reduce_sum(tf.square(W))
```

Dropout: A Simple Way to Prevent Neural Networks from Overfitting [Srivastava et al. 2014]




Regularization: Dropout

“randomly set some neurons to zero in the forward pass”





(a) Standard Neural Net



(b) After applying dropout.


[Srivastava et al., 2014]

Waaaait a second...
How could this possibly be a good idea?



TensorFlow implementation

```
dropout_rate = tf.placeholder("float")
_L1 = tf.nn.relu(tf.add(tf.matmul(X, W1), B1))
L1 = tf.nn.dropout(_L1, dropout_rate)
```




~~TRAIN:~~

```
sess.run(optimizer, feed_dict={X: batch_xs, Y: batch_ys,
dropout_rate: 0.7})
```


~~EVALUATION:~~

```
= print "Accuracy:", accuracy.eval({X: mnist.test.images, Y:
mnist.test.labels, dropout_rate: 1})
```

What is Ensemble?




Next
NN LEGO Play!




Lecture 10-4

NN LEGO Play



Sung Kim <hunkim+mr@gmail.com>


Feedforward neural network





Fast forward

He

3% ↓




Split & merge




Recurrent network

RNN



'The only limit is your imagination'



<http://itchyi.squarespace.com/thelatest/2012/5/17/the-only-limit-is-your-imagination.html>